

```
In [1]:  
  
import pandas as pd  
import numpy as np  
import seaborn as sns  
import matplotlib.pyplot as plt  
  
In [2]:  
  
data=pd.read_csv('health care diabetes.csv')
```

```
In [3]:  
  
data
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction
0	6	148	72	35	0	33.6	0.627
1	1	85	66	29	0	26.6	0.351
2	8	183	64	0	0	23.3	0.672
3	1	89	66	23	94	28.1	0.167
4	0	137	40	35	168	43.1	2.288
...
763	10	101	76	48	180	32.9	0.171
764	2	122	70	27	0	36.8	0.340
765	5	121	72	23	112	26.2	0.245
766	1	126	60	0	0	30.1	0.349
767	1	93	70	31	0	30.4	0.315

768 rows x 9 columns

```
In [4]:  
  
data.shape  
  
(768, 9)
```

In [5]:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   Pregnancies           768 non-null   int64  
 1   Glucose               768 non-null   int64  
 2   BloodPressure         768 non-null   int64  
 3   SkinThickness         768 non-null   int64  
 4   Insulin               768 non-null   int64  
 5   BMI                   768 non-null   float64 
 6   DiabetesPedigreeFunction 768 non-null   float64 
 7   Age                  768 non-null   int64  
 8   Outcome               768 non-null   int64  
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

No Null Values in Data

In [6]:

```
data['Pregnancies'].isnull().sum()
data['Glucose'].isnull().sum()
data['BloodPressure'].isnull().sum()
data['SkinThickness'].isnull().sum()
data['Insulin'].isnull().sum()
data['BMI'].isnull().sum()
data['DiabetesPedigreeFunction'].isnull().sum()
data['Age'].isnull().sum()
```

0

In [7]:

```
data.isnull().any()
```


```
Pregnancies      False
Glucose           False
BloodPressure     False
SkinThickness     False
Insulin           False
BMI               False
DiabetesPedigreeFunction  False
Age              False
Outcome           False
dtype: bool
```

In [8]:

```
data['Pregnancies'].value_counts()
```

```
plt.hist(data['Pregnancies'])
```

```
(array([246., 178., 125.,  50.,  83.,  52.,  11.,  19.,   3.,   1.]),
 array([ 0. ,  1.7,  3.4,  5.1,  6.8,  8.5, 10.2, 11.9, 13.6, 15.3, 17. ]),
 <BarContainer object of 10 artists>)
```



In [9]:

```
data['Pregnancies'].value_counts()
```


```
1    135
0    111
2    103
3     75
4     68
5     57
6     50
7     45
8     38
9     28
10    24
11    11
13    10
12     9
14     2
15     1
17     1
```

Name: Pregnancies, dtype: int64

In [10]:

```
plt.hist(data['Glucose'])
```


```
(array([ 5.,  0.,  4., 32., 156., 211., 163., 95., 56., 46.]),
 array([ 0., 19.9, 39.8, 59.7, 79.6, 99.5, 119.4, 139.3, 159.2,
        179.1, 199. ]),
 <BarContainer object of 10 artists>)
```



In [11]:

```
data['BloodPressure'].value_counts()  
plt.hist(data['BloodPressure'])
```


```
(array([ 35.,  1.,  2., 13., 107., 261., 243., 87., 14.,  5.]),  
array([ 0. , 12.2, 24.4, 36.6, 48.8, 61. , 73.2, 85.4, 97.6,  
       109.8, 122. ]),  
<BarContainer object of 10 artists>)
```



In [12]:

```
data['SkinThickness'].value_counts()  
plt.hist(data['SkinThickness'])
```


```
(array([231., 107., 165., 175., 78.,  9.,  2.,  0.,  0.,  1.]),  
array([ 0. ,  9.9, 19.8, 29.7, 39.6, 49.5, 59.4, 69.3, 79.2, 89.1, 99. ]),  
<BarContainer object of 10 artists>)
```



In [13]:

```
data['Insulin'].value_counts()  
plt.hist(data['Insulin'])
```


```
(array([487., 155., 70., 30., 8., 9., 5., 1., 2., 1.]),  
array([ 0. , 84.6, 169.2, 253.8, 338.4, 423. , 507.6, 592.2, 676.8,  
       761.4, 846. ]),  
<BarContainer object of 10 artists>)
```



In [14]:

```
data['BMI'].value_counts()  
plt.hist(data['BMI'])
```


```
(array([ 11.,  0., 15., 156., 268., 224., 78., 12.,  3.,  1.]),  
array([ 0. ,  6.71, 13.42, 20.13, 26.84, 33.55, 40.26, 46.97, 53.68,  
       60.39, 67.1 ]),  
<BarContainer object of 10 artists>)
```



In [15]:

```
data['DiabetesPedigreeFunction'].value_counts()
plt.hist(data['DiabetesPedigreeFunction'])
```

```
(array([318., 206., 136.,  58.,  25.,  15.,   3.,   3.,   1.,   3.]),
 array([0.078 , 0.3122, 0.5464, 0.7806, 1.0148, 1.249 , 1.4832, 1.7174,
        1.9516, 2.1858, 2.42  ]),
 <BarContainer object of 10 artists>)
```



In [16]:


```
data.describe().transpose()
```

	count	mean	std	min	25%	50%	75%	max
Pregnancies	768.0	3.845052	3.369578	0.000	1.00000	3.0000	6.00000	17.00
Glucose	768.0	120.894531	31.972618	0.000	99.00000	117.0000	140.25000	199.00
BloodPressure	768.0	69.105469	19.355807	0.000	62.00000	72.0000	80.00000	122.00
SkinThickness	768.0	20.536458	15.952218	0.000	0.00000	23.0000	32.00000	99.00
Insulin	768.0	79.799479	115.244002	0.000	0.00000	30.5000	127.25000	846.00
BMI	768.0	31.992578	7.884160	0.000	27.30000	32.0000	36.60000	67.10
DiabetesPedigreeFunction	768.0	0.471876	0.331329	0.078	0.24375	0.3725	0.62625	2.42
Age	768.0	33.240885	11.760232	21.000	24.00000	29.0000	41.00000	81.00
Outcome	768.0	0.348958	0.476951	0.000	0.00000	0.0000	1.00000	1.00

In [17]:

```
sns.scatterplot(x='Pregnancies',y='Glucose',hue='Outcome',data=data)
```

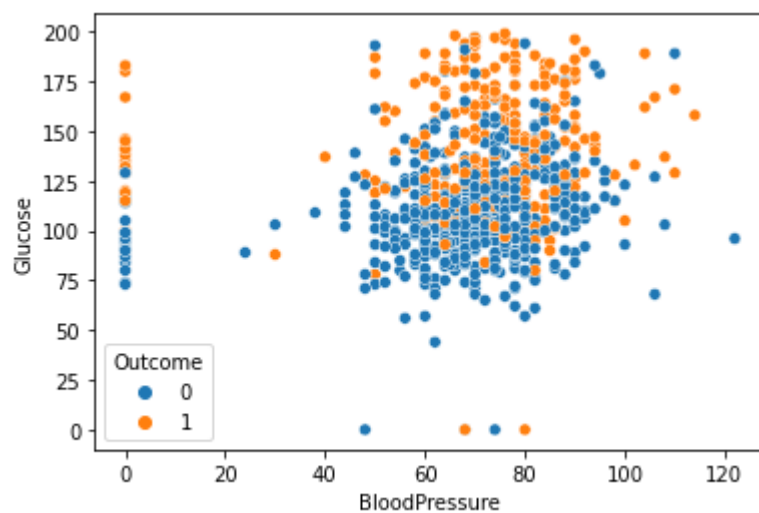
<AxesSubplot:xlabel='Pregnancies', ylabel='Glucose'>



In [18]:

```
sns.scatterplot(x='BloodPressure',y='Glucose',hue='Outcome',data=data)
```


<AxesSubplot:xlabel='BloodPressure', ylabel='Glucose'>



In [19]:

```
sns.scatterplot(x='SkinThickness',y='Glucose',hue='Outcome',data=data)
```


<AxesSubplot:xlabel='SkinThickness', ylabel='Glucose'>



In [20]:

```
sns.scatterplot(x='BMI',y='Insulin',hue='Outcome',data=data)
```


<AxesSubplot:xlabel='BMI', ylabel='Insulin'>



In [21]:

```
sns.scatterplot(x='Age',y='Insulin',hue='Outcome',data=data)
```

<AxesSubplot:xlabel='Age', ylabel='Insulin'>



In [22]:


```
data.corr()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
Pregnancies	1.000000	0.129459	0.141282	-0.081672	-0.073535	0.017683	-0.033258	0.544341	0.221898
Glucose	0.129459	1.000000	0.152590	0.057328	0.331357	0.221071	0.137337	0.263514	0.466581
BloodPressure	0.141282	0.152590	1.000000	0.207371	0.088933	0.281805	0.041265	0.239528	0.065068
SkinThickness	-0.081672	0.057328	0.207371	1.000000	0.436783	0.392573	0.183928	-0.113970	0.074752
Insulin	-0.073535	0.331357	0.088933	0.436783	1.000000	0.197859	0.185071	-0.042163	0.130548
BMI	0.017683	0.221071	0.281805	0.392573	0.197859	1.000000	0.140647	0.036242	0.292695
DiabetesPedigreeFunction	-0.033258	0.137337	0.041265	0.183928	0.185071	0.140647	1.000000	0.258160	0.672867
Age	0.544341	0.263514	0.239528	-0.113970	-0.042163	0.036242	0.258160	1.000000	0.082351
Outcome	0.221898	0.466581	0.065068	0.074752	0.130548	0.292695	0.672867	0.082351	1.000000

In [23]:

```
plt.figure(figsize=(10,7))
sns.heatmap(data.corr(),annot=True,cmap="YlGnBu")
```

<AxesSubplot:>



Model Building

In [24]:

```
X=data.iloc[:,[0,1,2,3,4,5,6,7]]
```

In [25]:

```
y=data.iloc[:,[8]]
```

In [26]:

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_stat
```

In [27]:

```
from sklearn.linear_model import LogisticRegression
```

In [28]:

```
model= LogisticRegression()
model.fit(X_train,y_train)
```

d:\Users\coold\anaconda3\lib\site-packages\sklearn\utils\validation.py:993: DataConversionWarning: A column-\narray was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
y = column_or_1d(y, warn=True)
```

d:\Users\coold\anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:814: ConvergenceWarning: lbfgs f\nSTOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.ht

```
n_iter_i = _check_optimize_result(
```

```
LogisticRegression()
```

In [29]:

```
print(model.score(X_train,y_train))
print(model.score(X_test,y_test))
```

```
0.7801556420233463
```

```
0.7480314960629921
```

In [30]:

```
from sklearn.metrics import confusion_matrix,accuracy_score
from sklearn.metrics import classification_report
```

In [31]:

```
yhat=model.predict(X_test)
```

In [32]:

```
print(accuracy_score(y_test,yhat))
```

```
0.7480314960629921
```

In [33]:

```
print(confusion_matrix(y_test,yhat))
```

```
[[136  32]
 [ 32  54]]
```

In [34]:

```
print(classification_report(y_test,yhat))
```

	precision	recall	f1-score	support
0	0.81	0.81	0.81	168
1	0.63	0.63	0.63	86
accuracy			0.75	254
macro avg	0.72	0.72	0.72	254
weighted avg	0.75	0.75	0.75	254

In [35]:

```
from sklearn.neighbors import KNeighborsClassifier
model2 = KNeighborsClassifier(p=2,n_neighbors=5,metric='minkowski')
```

In [36]:

```
model2.fit(X_train,y_train)
```

```
d:\Users\coold\anaconda3\lib\site-packages\sklearn\neighbors\_classification.py:198: DataConversionWarning: A
hen a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
    return self._fit(X, y)
```

```
KNeighborsClassifier()
```

In [37]:

```
print(model2.score(X_train,y_train))
print(model2.score(X_test,y_test))
yhat2=model2.predict(X_test)
print(confusion_matrix(y_test,yhat2))
print(accuracy_score(y_test,yhat2))
```

```
0.7937743190661478
0.7007874015748031
[[130  38]
 [ 38  48]]
0.7007874015748031
```

In [38]:

```
from sklearn.ensemble import RandomForestClassifier
model3=RandomForestClassifier(n_estimators=15,criterion='gini',max_features='auto')
```

In [39]:

```
model3.fit(X_train,y_train)
print(model3.score(X_train,y_train))
print(model3.score(X_test,y_test))
```

```
0.9922178988326849
0.7401574803149606
```

C:\Users\coold\AppData\Local\Temp\ipykernel_18212\2588249225.py:1: DataConversionWarning: A column-vector y was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
model3.fit(X_train,y_train)
```

In [40]:

```
from sklearn.metrics import roc_curve,roc_auc_score
```

In [44]:

```
p1=model.predict_proba(X_train)
p2=model2.predict_proba(X_train)
p3=model3.predict_proba(X_train)
```

```
In [53]:  
  
## keeping probs for positive outcomes only  
p1=p1[:,1]  
p2=p2[:,1]  
p3=p3[:,1]
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```